

Spark streaming

~~Exactly once~~

At least once + idempotence

Quentin Ambard
@qambard


Agenda

- Native checkpointing
 - Driver & executor failure
- Custom checkpointing
- Performance tips
- Structured streaming (spark 2.2)
- Logs
- Kafka
- Bonus: Kafka 1 & “exactly-once”

Spark & kafka for (near)realtime



Once up a time, Spark met Kafka and...



Spark native checkpointing



Native checkpointing

Spark saves for you the job state to a resilient storage (dsefs) :

- Save metadata, serialize the streaming context with:
 - Spark configuration
 - DStream operations (what your job is doing)
 - Incomplete batches (scheduled but not yet executed)
 - Last kafka offsets

Checkpoints twice: before and after each micro-batch (blocking)



Enabled with `ssc.checkpoint(dir)`

- Save the actual RDD data.
Should be enabled for stateful streams only (sliding window etc)
Expensive, enabled with `kafkaStream.checkpoint(Seconds(10))`

Checkpointing: example

```
val dir = "dsefs://35.188.108.191:5598/exactlyonce"
```

```
def createStreamingCtx() = {  
  val ssc = new StreamingContext(conf, Seconds(5))  
  val stream = KafkaUtils.createDirectStream(ssc, kafkaParams, topics)  
  //Apply transformations on each microbatch  
  stream.foreachRDD(rdd => println(rdd.count()))  
  //Set checkpointing directory  
  ssc.checkpoint(dir)  
  ssc  
}  
val ssc = StreamingContext.getOrCreate(dir, createStreamingCtx)
```



Checkpointing: restart behavior

- Read data saved in the checkpointing directory
- Schedule all missing job:
1 hour outage with a 1 seconds window will generate 60*60 micro-batches (!).
- First micro-batch starts from the last saved offsets
- First micro-batch will read all available kafka messages
Can be limited with
 - `spark.streaming.kafka.maxRatePerPartition = XXX`
 - `spark.streaming.backpressure.enabled = true`

Checkpointing: serialization

Serialize Dstream to disk

- Need to be serializable
- Won't be able to restart a job if the Dstream operations have changed (classes are different)

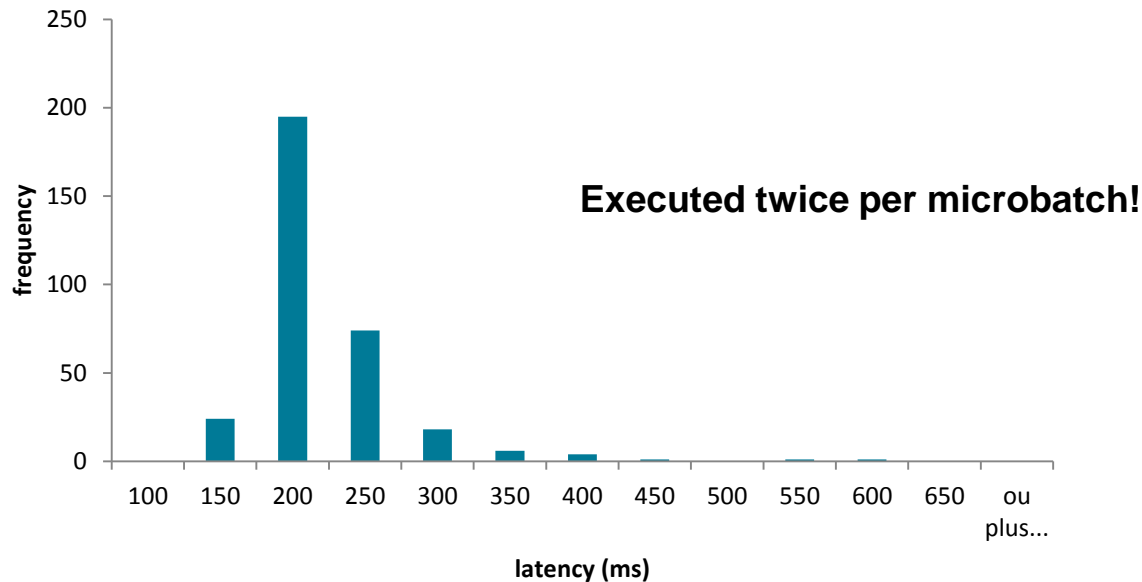
This problem alone discards native checkpointing for a lot of customers

Workaround:

- Run the job twice in parallel during upgrade (can be an issue if the order matters)
- “Clean stop”, restart from the beginning of the kafka topic

Checkpointing: speed

Basic job, saved to DSEFS (7.7 kb), 6 nodes




Failure handling



Worker failure

Scenario: the worker running the driver get killed



- Master will restart the driver on another worker
- Driver will stop... Sometime... (watches its worker)
- 2 spark apps running at the same time for a few sec
- Can break order
- Same checkpointing dir

Worker failure

Scenario: worker running the driver get killed

Concurrent operation on dsefs?



Worker failure

Workaround:

Make sure it's the only instance running before starting the stream

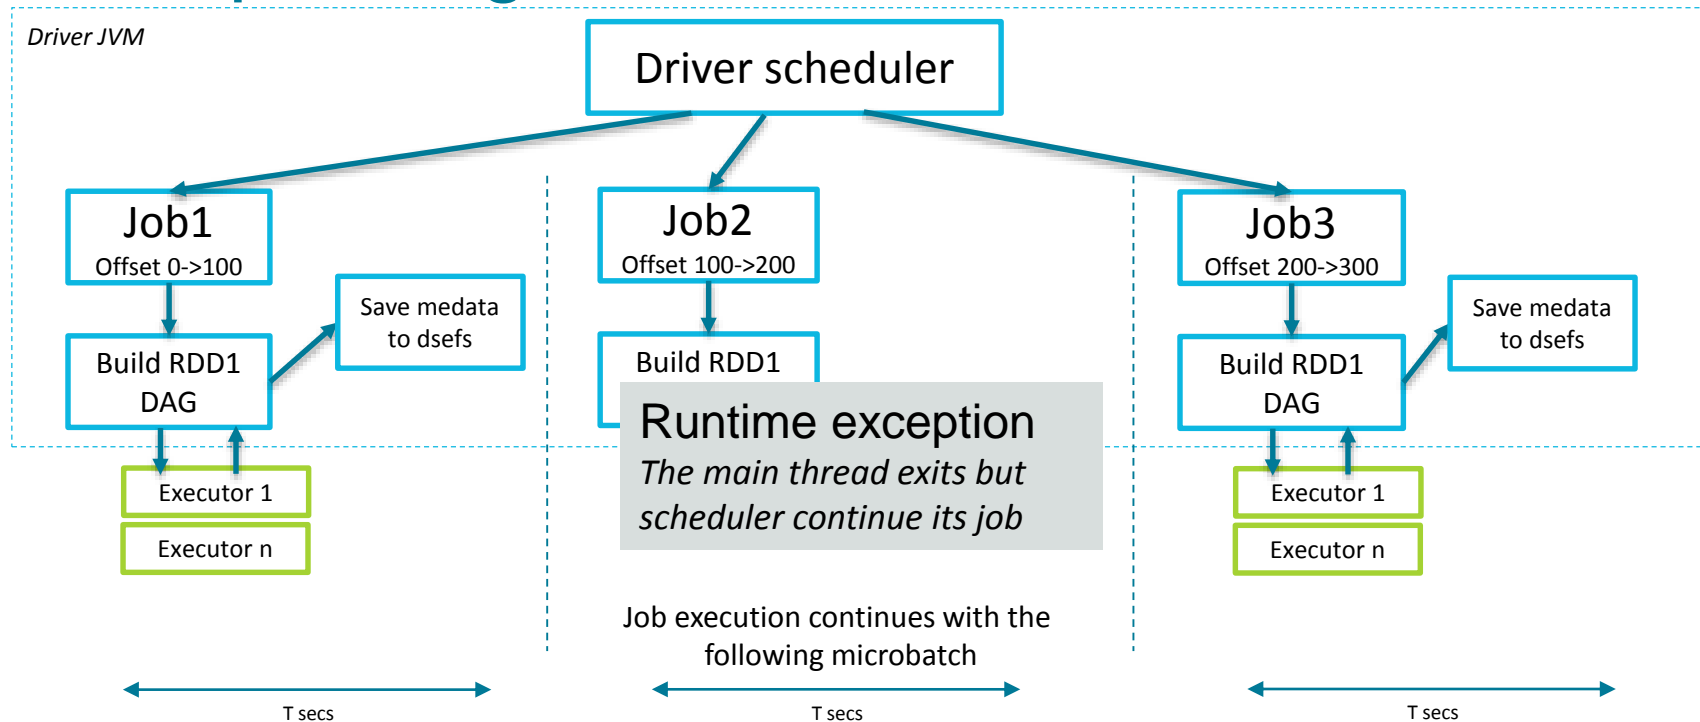
```
curl http://35.188.90.197:7080/api/v1/applications
```

or curl secret api ?

Checkpointing: failure in the driver

```
kafkaStream.foreachRDD(rdd => {  
    i += 1  
    println(s"Processing RDD $i")  
    if (i == 3) {  
        throw new RuntimeException(s"Exception for RDD $i")  
    }  
    val count = rdd.map(_._2+ "computed ").count()  
    println(s"RDD $i completed : $count")  
})
```

Checkpointing: failure in the driver



Data from Job #2 (offset 100->200) is lost and won't be recovered (erased by the next checkpointing)

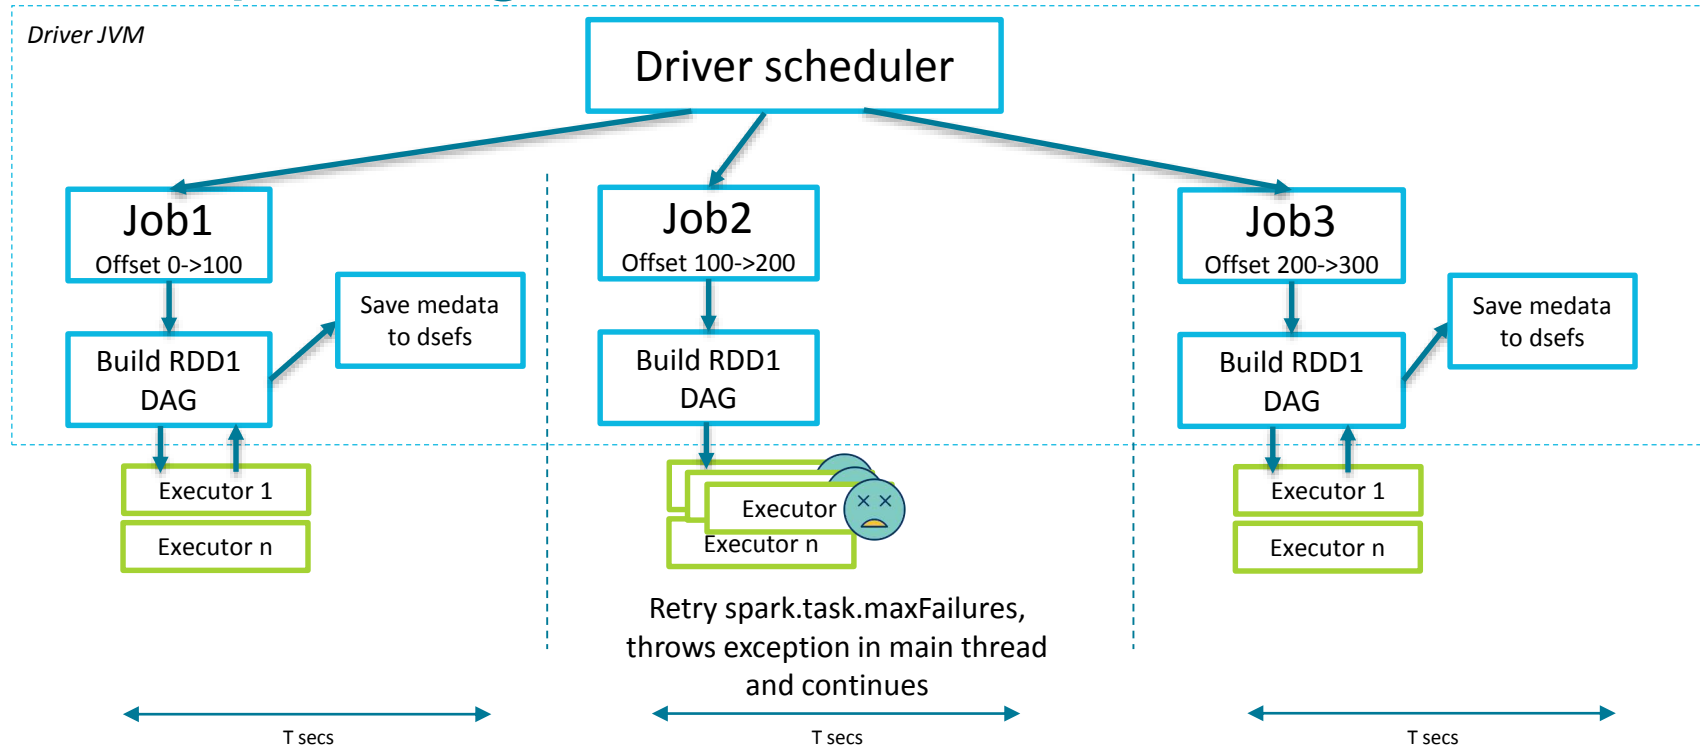
Checkpointing: failure in an executor

Inside
an executor



```
kafkaStream.foreachRDD(rdd => {  
  i += 1  
  println(s"Processing RDD $i")  
  val count = rdd.map(row => {  
    {  
      row._2+ "computed "  
      if (i == 3) {  
        throw new RuntimeException(s"Exception for RDD $i")  
      }  
    }).count()  
  })  
  println(s"RDD $i completed : $count")  
})
```

Checkpointing: failure in the driver



Data from Job #2 (offset 100->200) is lost and won't be recovered (erased by the next checkpointing)

Checkpointing: fail-fast

Need a solution to stop job execution as soon as a micro-batch fail


<https://issues.apache.org/jira/browse/SPARK-6415>

Workaround:

- Catch the error and kill the driver asap ?

Checkpointing: fail-fast

```
ssc.start()
Try {
  ssc.awaitTermination()
} match {
  case Success(result) => println("StreamingContext has been stopped")
  case Failure(exception) => println("Stopping StreamingContext driver...", exception)
  ssc.stop(stopSparkContext = true, stopGracefully = false)
  System.exit(50) //sparkExitCode.UNCAUGHT_EXCEPTION
}
```



Checkpointing: fail-fast

- Streaming context waits 2 seconds before stopping (with `graceful = false`).
- Scheduler continues its job and launch the next batch
- Shutdown hook to stop context on `System.exit()`
- Managed to reproduce with a `Runtime.getRuntime.halt()`
- Not guarantee that the next microbatch won't be checkpointed

Conclusion

A close-up, black and white photograph of a pug's face. The dog's eyes are large and dark, looking directly at the camera. Its snout is prominent in the center, with visible nostrils. The skin is heavily wrinkled, especially around the eyes and on the snout. The ears are large and floppy, visible on the sides of the head. The overall texture of the fur appears soft but with many fine wrinkles.

Checkpointing: conclusion



Easy to setup (yeee)



Reschedule all missed jobs during startup

Will likely kill your driver if it was stopped for a few days :/



Code/spark upgrade non-trivial (serialization issue)



Can't easily monitor Kafka offset consumption



Slow



Lot of corner-cases



Couldn't find a way to make it 100% reliable

Custom checkpointing



Read kafka offset from rdd

Stored with RDD metadata

```
val offsetRanges = rdd.asInstanceOf[HasOffsetRanges].offsetRanges

offsetRanges.foreach{range =>
    println(s"OFFSET RANGES for partition :${range.partition} =
    [${range.fromOffset}-${range.untilOffset}])
}
//OFFSET RANGES for partition :1 = [0-10]
//OFFSET RANGES for partition :2 = [0-13]
...
```

Store kafka offset

Store data in C* yourself instead of hdfs

```
def saveOffsetRanges (topic, ranges) = {  
  val rangesJson = ranges.sortBy(_.partition).toJson()  
  session.execute("insert into ks.checkpoint (topic, offsets) values (  
    ${topic}, ${rangesJson} ")  
}
```

Doesn't handle failures better than native checkpointing

Need to make sure the next batch won't checkpoint until the previous is actually saved

Store kafka offset

Use a LWT to guarantee exactly-once

```
def saveOffsetRanges (topic, ranges, previousRanges) = {  
  val rangesJson = ranges.sortBy(_.partition).toJson()  
  val rs = session.execute(" update ks.checkpoint set offsets=  
    ${rangesJson} where topic=${topic} IF offsets=${previousRanges} ")  
  if (!rs.wasApplied){  
    log.error("last offsets doesn't match. Killing driver. ")  
    System.exit(50)  
  }  
  rangesJson  
}
```

Ensure application uniqueness

//generated and updated to ks.checkpoint at application startup

```
val appUuid = UUIDs.random()
```

```
session.execute("update ks.checkpoint set app_id= ${appUuid} where  
topic=${topic}")
```


....

```
def saveOffsetRanges (topic, ranges, previousRanges) = {  
    val rangesJson = ranges.sortBy(_.partition).toJson()  
    val rs = session.execute("update ks.checkpoint set offsets=  
${rangesJson} where topic=${topic} IF offsets=${previousRanges} and  
app_id=${appUuid} ")  
    if (!rs.wasApplied){  
        ...  
    }  
}
```


When should we save offsets ?

Always after the executor work

```
kafkaStream.transform(rdd => {  
    rdd.doSomething...  
}).foreachRDD(rdd => {  
    rdd.doSomething(...)  
    val offsetRanges = rdd.asInstanceOf[HasOffsetRanges].offsetRanges  
    previousRanges = saveOffsetRanges(offsetRanges)  
})
```




If you need to save the state at the beginning



```
kafkaStream.transform(rdd => {  
    //executed by job scheduler (LWT will fail with the first lag)  
    saveInitialState(rdd, topic, previousRange)  
}).foreachRDD(rdd => {  
    rdd.doSomething...  
    previousRanges = saveOffsetRanges (rdd, topic, previousRange)  
})
```

If you need to save the state at the beginning



```
kafkaStream.foreachRDD(rdd => {  
    //Must be protected by LWT  
    saveInitialState(rdd, topic, previousRange)  
    rdd.doSomething...  
    previousRanges = saveOffsetRanges (rdd, topic, previousRange)  
})
```

Custom checkpointing: conclusion

- ✓ Reliable
- ✓ Kafka offset consumption easy to monitor
- ✓ Easy upgrade
- ✓ Fast
- ✗ Harder to setup (Need to use LWT)

Performances tips



Performances

Very easy to get dead-slow throughput with complex jobs

(15 messages/sec for a solid 3 nodes cluster!)

The usual recommendations apply:

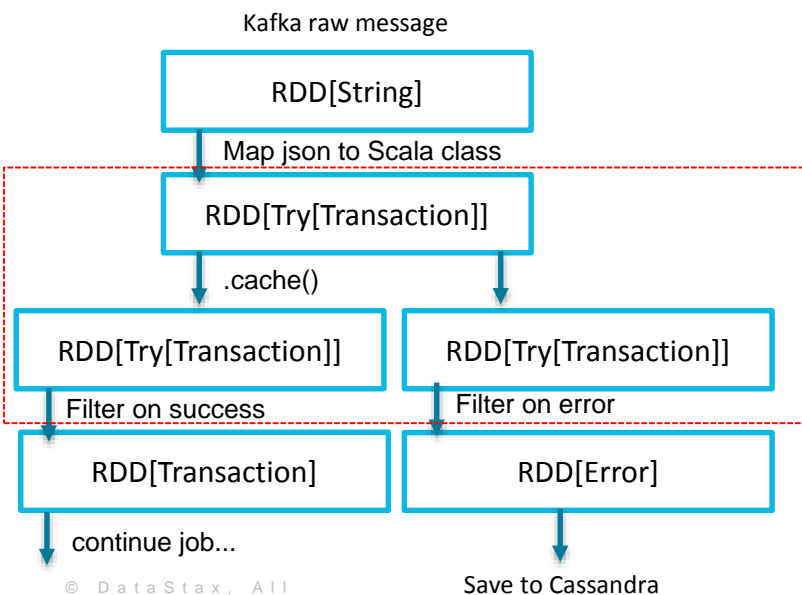
- Data distribution
- Avoid shuffle
- Prefer DataSet
- Use reducer instead of groupBy when possible...

Many kafka partitions with small process time?

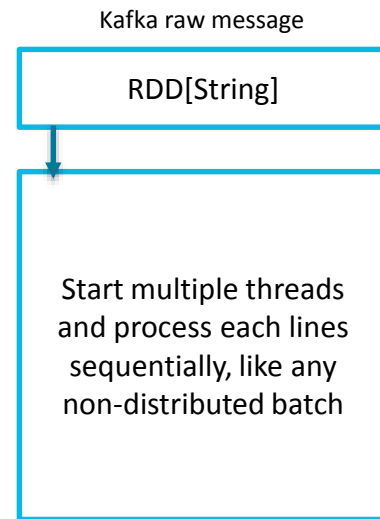
- `rdd.coalesce(numberOfExecutor)`
- Disable `spark.locality.wait`

Performances

- Doing a lot of `.cache()` with multiple actions will hurt
- It's often easier and faster to parallelize synchronous computation in your executors



200ms: fast for long
batch, slow for a 3
seconds micro-batch



<http://www.russellspitzer.com/2017/02/27/Concurrency-In-Spark/>

Performances

Monitor the driver and executor JVM !!

Typical scenario on driver:

- microbatches get slow for some reason
- Scheduler schedules many batches, add pressure on the heap
- Many GC, even slower, more lag, more GC, more lag, more GC...

On executors:

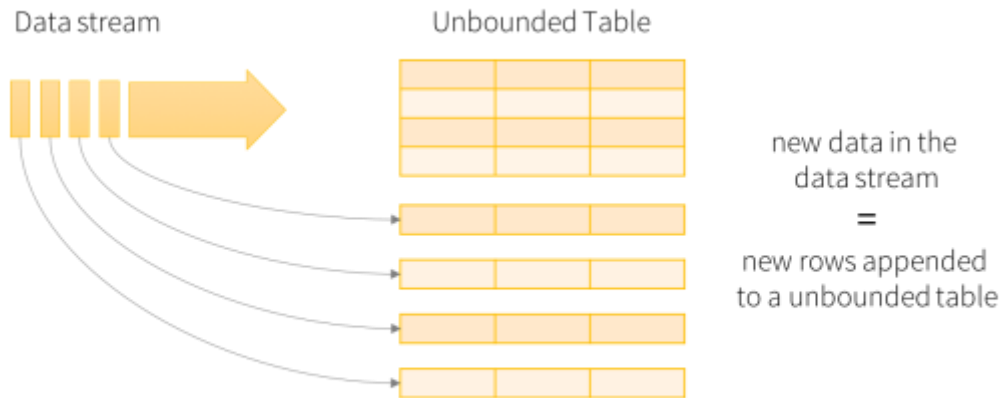
- By default, heap at 1GB
- “Only” $(1024\text{MB} - 300\text{MB}) * 0.25 = 180\text{MB}$ for your objects!
(the rest is for spark storage and execution)

Structured streaming (2.2)



Unified API

- Uses DataSet instead of RDD
- Better support for aggregations
- Based on message timestamp
- Support “late data” (watermark)
- Stop/start queries
- ...



Data stream as an unbounded table

Example

```
val ds = sparkSession.readStream
  .format("kafka")
  .option("kafka.bootstrap.servers", "localhost:9092")
  .option("subscribe", "myTopic")
  .load()
  .select(from_json($"value", schema).as("json"))
  .select("json.*").as[Transaction]
```

```
val query = ds.writeStream
  .cassandraFormat("transaction", ks)
  .outputMode(OutputMode.Update)
  .start()
```

How checkpointing works

- Checkpoint data to resilient storage (DSEFS, HDFS etc.)
- Once after each microbatch
- Conceptually similar to custom checkpointing
- Data saved as JSON

```
{"batchWatermarkMs":0,"batchTimestampMs":1502995346109,"conf":{"spark.sql.shuffle.partitions":"200"}}  
{"exactlyonce":{"0":32625}}
```

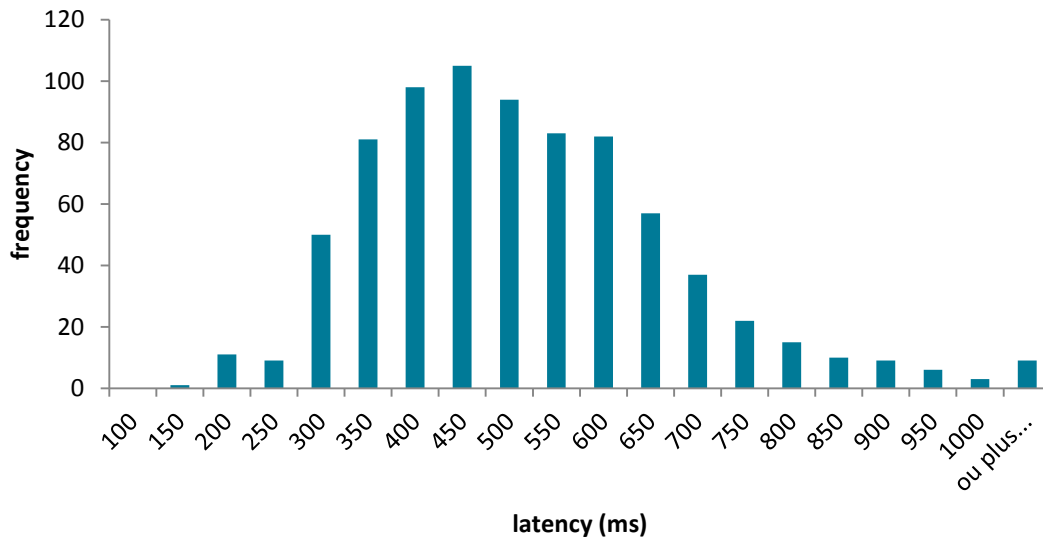

Failures

- Fails fast and safely!
- Catch exception and safely restart the query without killing the driver
- Still need to ensure the same query isn't running twice

Checkpointing speed

Checkpointing latency

spark 2.2 -> DSEFS 5.1.2



DSEFS operations per microbatch: 9 read, 2 creation, 2 rename, 2 deletes

Speed & Triggers

- Equivalent to window size
- 0 by default (= proceed asap)

```
ds.writeStream  
  .trigger(Trigger.ProcessingTime("1 seconds"))  
  .queryName("exactlyOnce").foreach(writer).start
```

- Increase throughput with a higher maxOffsetsPerTrigger

timeToProcess(maxOffsetsPerTrigger) >> checkpointing time
(otherwise spend too much time checkpointing)

Faster checkpointing?

Implement a custom FileSystem to store on C* only (instead of DSEFS)

```
class CassandraSimpleFileSystem extends hadoop.fs.FileSystem {  
  override def mkdirs(f: Path, permission: FsPermission): Boolean = ???  
  ...  
}
```

```
SparkSession.builder.config("spark.hadoop.fs.ckfs.impl",  
  "exactlyonce.CassandraSimpleFileSystem")
```


```
ds.writeStream.option("checkpointLocation",  
  "ckfs://127.0.0.1/checkpointing/exactlyOnce")  
  .queryName("exactlyOnce").foreach(writer).start
```

<https://github.com/QuentinAmbard/cassandracheckointingfs/>

Custom checkpointing

Checkpointing latency

spark 2.2 -> C* custom FS



Structured streaming: conclusion

- ✓ Reliable
- ✓ Fast
- ✓ DataSet & powerful time-based aggregation (support for late data)
- ✓ Easy upgrade (as long as aggregation type doesn't change)
- ✓ It's the future!
- ✗ Currently slow with DSEFS
- ✗ Old (blocked) data could be discarded if using producer time aggregation.
Use custom sink to reproduce previous streaming behavior



Logs

Logs

spark-env.sh

```
export SPARK_WORKER_OPTS="$SPARK_WORKER_OPTS -  
Dspark.worker.cleanup.enabled=true -  
Dspark.worker.cleanup.interval=300 -  
Dspark.worker.cleanup.appDataTtl=300"
```

Cleanup non-running application only

Won't help for long-living streaming job

Logs

Limit the executor log size:

(SPARK_WORKER_DIR/worker-n/application_id/executor_id/...)

spark-defaults.conf

```
spark.executor.logs.rolling.maxRetainedFiles 3
```

```
spark.executor.logs.rolling.strategy size
```

```
spark.executor.logs.rolling.maxSize 50000
```

Logs

- If an executor jvm crashes, a new executor is started. Creates new folder with a copy of the application .jar
- Crashed executors folder won't be cleanup.
- Good practice to monitor the size of spark worker directory

(but your executors shouldn't OOM too often right?)

Kafka



Kafka durability

- Disable unclean leader election (now default)
- Define in-synch / RF (ex: RF=3, insynch=2)
- Publish at ALL (ack from all insynch)
- If order matters, allow only 1 inflight request (or 0 retries)

```
unclean.leader.election.enable = false
```

```
default.replication.factor = 3
```

```
min.insync.replicas=2
```

```
#producer
```

```
max.in.flight.requests.per.connection = 1
```

```
retries=3
```

```
acks = all
```

```
#consumer (disable autocommit for direct stream)
```

```
enable.auto.commit = false
```

How many kafka partitions ?

- Lot of partition will slow things down
- ~3x your number of spark cores is a good start
- Must allow to scale if needed
 - If order matters, changing changing kafka partitioning can be hard

Kafka 11: idempotent producer

- Idempotent producer:
 - Generate a unique id **per connection/producer**
Your executors might restart with different producers id!
 - Retry an unlimited number of time
 - Broker de-duplicate based on the message id

Kafka 11: Transactions

```
props.put("transactional.id", "myId")  
...  
producer.initTransactions()  
  
producer.beginTransaction()  
kafka.producer.send(myProducerRecord1)  
kafka.producer.send(myProducerRecord2)  
kafka.producer.commitTransaction()
```

```
//Consumer configuration  
props.put("isolation.level", "read_committed")
```

Kafka 11: Transactions

- Consumer blocked to the last stable offset (LSO)
 - 1 transaction stalled => LSO blocked => consumer blocked
- `Transaction.max.timeout.ms` set to 900000 (15min) by default
- “Exactly-once streams” based on transactions

Thank you